

Detecting Fake News Using Sentiment Analysis

Andrea Boskovic and Peter Cho

12/18/2018

```
# Loading the dataset
fake <- read_csv("fake.csv") #all fake

## Parsed with column specification:
## cols(
##   .default = col_character(),
##   ord_in_thread = col_integer(),
##   published = col_datetime(format = ""),
##   crawled = col_datetime(format = ""),
##   domain_rank = col_integer(),
##   spam_score = col_double(),
##   replies_count = col_integer(),
##   participants_count = col_integer(),
##   likes = col_integer(),
##   comments = col_integer(),
##   shares = col_integer()
## )

## See spec(...) for full column specifications.

real <- read_csv("Articles.csv") #all real

## Parsed with column specification:
## cols(
##   Article = col_character(),
##   Date = col_character(),
##   Heading = col_character(),
##   NewsType = col_character()
## )

new_ds <- read_csv("data.csv") #combination of real and fake

## Parsed with column specification:
## cols(
##   URLs = col_character(),
##   Headline = col_character(),
##   Body = col_character(),
##   Label = col_integer()
## )

fake_type <- c("fake", "satire", "bias", "bs", "conspiracy", "state", "junksci", "hate")
real_type <- c("sports", "business")

# Merging the datasets and removing unnecessary columns
real <- real %>%
  mutate(binary_type = ifelse(NewsType %in% fake_type, 0, 1)) #now fake = 0 and real = 1
fake <- fake %>%
  mutate(binary_type = ifelse(type %in% fake_type, 0, 1)) #now fake = 0 and real = 1
new_ds <- new_ds %>%
  filter(Label == 1)
```

```

real <- full_join(real, new_ds, by = c("Heading" = "Headline", "Article" = "Body", "binary_type" = "Label"))
real <- real %>%
  mutate(id = as.character(seq(1:4564))) %>%
  mutate(realtype = "real")

# Making a combined dataset with both fake and real articles and selecting only for the uuid (unique id)
combined <- full_join(fake, real, by = c("text" = "Article", "title" = "Heading", "uuid" = "id", "binary_type" = "Label"))
select(uuid, binary_type, type, title, text)

# Making a tidy dataset where we have the the words in their own column for facilitated data analysis
tidy_combined <- combined %>%
  unnest_tokens(word, text)

# Basic Data Exploration:

# This allows us to see how many observations are in each type of fake news.
combined %>%
  group_by(type) %>%
  summarize(n = n())

## # A tibble: 9 x 2
##   type      n
##   <chr>    <int>
## 1 bias      443
## 2 bs     11492
## 3 conspiracy  430
## 4 fake       19
## 5 hate      246
## 6 junksci   102
## 7 real     4564
## 8 satire    146
## 9 state     121

typetotals <- combined %>%
  group_by(type) %>%
  summarize(n = n())

# What are the most common words for each basic emotion?
# We will use the nrc lexicon to categorize each documented word into one of the basic human emotions categories

# Anger
nrc_anger <- get_sentiments("nrc") %>%
  filter(sentiment == "anger")

tidy_combined %>%
  inner_join(nrc_anger) %>%
  count(word, sort = TRUE)

## Joining, by = "word"

## # A tibble: 1,220 x 2
##   word      n
##   <chr>    <int>
## 1 vote     4969
## 2 money    4835
## 3 force    3189

```

```
## 4 court      2721
## 5 attack     2548
## 6 defense    2242
## 7 death      2176
## 8 bad        2175
## 9 politics   2058
## 10 fight     2054
## # ... with 1,210 more rows
```

```
# Fear
nrc_fear <- get_sentiments("nrc") %>%
  filter(sentiment == "fear")

tidy_combined %>%
  inner_join(nrc_fear) %>%
  count(word, sort = TRUE)
```

```
## Joining, by = "word"
## # A tibble: 1,430 x 2
##   word      n
##   <chr>    <int>
## 1 government 11656
## 2 war        9845
## 3 military   5880
## 4 police     4902
## 5 change     4442
## 6 case       4177
## 7 force      3189
## 8 court      2721
## 9 attack     2548
## 10 problem   2381
## # ... with 1,420 more rows
```

```
# Anticipation
nrc_anticipation <- get_sentiments("nrc") %>%
  filter(sentiment == "anticipation")

tidy_combined %>%
  inner_join(nrc_anticipation) %>%
  count(word, sort = TRUE)
```

```
## Joining, by = "word"
## # A tibble: 816 x 2
##   word      n
##   <chr>    <int>
## 1 time     14159
## 2 white     6547
## 3 public    6039
## 4 good      5802
## 5 long      5706
## 6 vote      4969
## 7 money     4835
## 8 investigation 3968
## 9 top       3822
## 10 continue  3439
```

```
## # ... with 806 more rows

# Trust
nrc_trust <- get_sentiments("nrc") %>%
  filter(sentiment == "trust")

tidy_combined %>%
  inner_join(nrc_trust) %>%
  count(word, sort = TRUE)

## Joining, by = "word"

## # A tibble: 1,191 x 2
##   word      n
##   <chr>   <int>
## 1 president 12344
## 2 united    7803
## 3 white     6547
## 4 good      5802
## 5 law       5181
## 6 system    5088
## 7 vote      4969
## 8 police    4902
## 9 money     4835
## 10 fact     4673
## # ... with 1,181 more rows

# Surprise
nrc_surprise <- get_sentiments("nrc") %>%
  filter(sentiment == "surprise")

tidy_combined %>%
  inner_join(nrc_surprise) %>%
  count(word, sort = TRUE)

## Joining, by = "word"

## # A tibble: 518 x 2
##   word      n
##   <chr> <int>
## 1 trump 23953
## 2 good  5802
## 3 vote  4969
## 4 money 4835
## 5 deal  2802
## 6 death 2176
## 7 leave 2080
## 8 hope  1902
## 9 young 1859
## 10 shot 1604
## # ... with 508 more rows

# Sadness
nrc_sadness <- get_sentiments("nrc") %>%
  filter(sentiment == "sadness")

tidy_combined %>%
```

```

inner_join(nrc_sadness) %>%
count(word, sort = TRUE)

## Joining, by = "word"

## # A tibble: 1,151 x 2
##   word      n
##   <chr>   <int>
## 1 vote    4969
## 2 black   4196
## 3 case    4177
## 4 problem 2381
## 5 lost    2260
## 6 tax     2211
## 7 death   2176
## 8 bad     2175
## 9 leave   2080
## 10 violence 1955
## # ... with 1,141 more rows

# Joy
nrc_joy <- get_sentiments("nrc") %>%
  filter(sentiment == "joy")

tidy_combined %>%
  inner_join(nrc_joy) %>%
  count(word, sort = TRUE)

## Joining, by = "word"

## # A tibble: 668 x 2
##   word      n
##   <chr> <int>
## 1 white  6547
## 2 good   5802
## 3 vote   4969
## 4 money  4835
## 5 found  4192
## 6 share  3090
## 7 deal   2802
## 8 food   2756
## 9 pay    2339
## 10 true   2234
## # ... with 658 more rows

# Disgust
nrc_disgust <- get_sentiments("nrc") %>%
  filter(sentiment == "disgust")

tidy_combined %>%
  inner_join(nrc_disgust) %>%
  count(word, sort = TRUE)

## Joining, by = "word"

## # A tibble: 1,023 x 2
##   word      n

```

```
##      <chr>      <int>
## 1 john          3108
## 2 congress      2473
## 3 death          2176
## 4 bad            2175
## 5 criminal       1805
## 6 illegal        1756
## 7 powerful       1611
## 8 corruption     1571
## 9 finally        1442
## 10 remains       1244
## # ... with 1,013 more rows
```

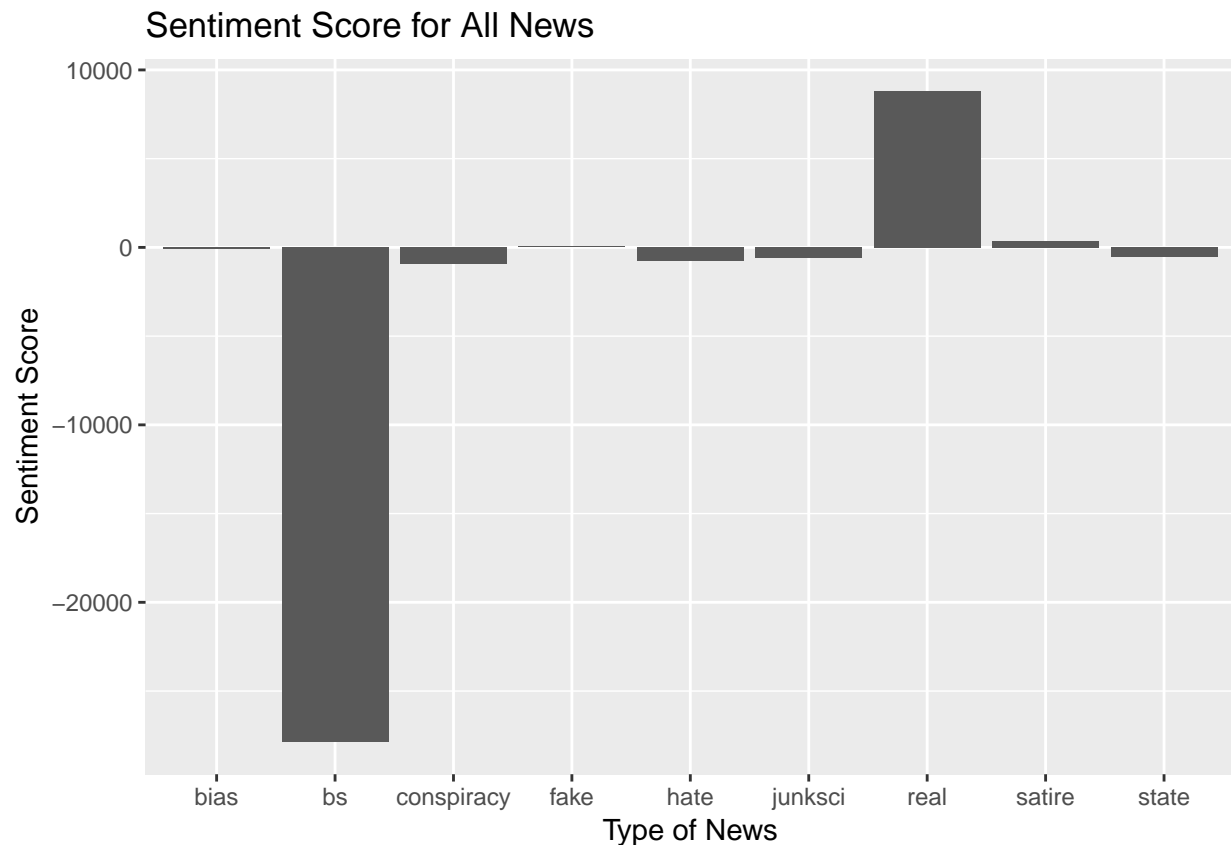
```
# Find net sentiment for each type of fake news documented in the dataset using the bing lexicon. The b
# Note that some types, such as bs (> 400,000), have more corresponding observations than other types, .
combined_sentiment <- tidy_combined %>%
  inner_join(get_sentiments("bing")) %>%
  count(type, sentiment) %>%
  spread(sentiment, n, fill = 0) %>%
  mutate(sentiment = positive - negative)
```

```
## Joining, by = "word"
```

```
combined_sentiment
```

```
## # A tibble: 9 x 4
##   type      negative positive sentiment
##   <chr>      <dbl>    <dbl>    <dbl>
## 1 bias          5422     5322     -100
## 2 bs        247391    219536   -27855
## 3 conspiracy   4805      3851    -954
## 4 fake         148       199       51
## 5 hate         8765     7998    -767
## 6 junksci      3070     2469    -601
## 7 real        45896    54690    8794
## 8 satire       1148     1487     339
## 9 state        1215      704    -511
```

```
# Plot of the sentiment score for each type of news
ggplot(combined_sentiment, aes(x = type, y = sentiment)) + geom_col() + labs(title = "Sentiment Score f
```



```
# We can also get the sentiment score on a scale of -5 to 5 from the AFINN lexicon. The AFINN lexicon h
afinn <- tidy_combined %>%
  inner_join(get_sentiments("afinn")) %>%
  group_by(type) %>%
  summarise(sentiment = sum(score)) %>%
  mutate(method = "AFINN")
```

```
## Joining, by = "word"
```

```
afinn
```

```
## # A tibble: 9 x 3
##   type      sentiment method
##   <chr>      <int> <chr>
## 1 bias      -1507 AFINN
## 2 bs       -6201 AFINN
## 3 conspiracy -1846 AFINN
## 4 fake        108 AFINN
## 5 hate     -1625 AFINN
## 6 junksci     41 AFINN
## 7 real      28457 AFINN
## 8 satire      868 AFINN
## 9 state     -1089 AFINN
```

```
# It may be useful to investigate the basic contents of the lexicons.
```

```
# Positive and negative words in nrc lexicon
```

```
get_sentiments("nrc") %>%
  filter(sentiment %in% c("positive",
```

```

    "negative")) %>%
count(sentiment)

## # A tibble: 2 x 2
##   sentiment      n
##   <chr>      <int>
## 1 negative   3324
## 2 positive   2312

# Positive and negative words in bing lexicon
get_sentiments("bing") %>%
count(sentiment)

## # A tibble: 2 x 2
##   sentiment      n
##   <chr>      <int>
## 1 negative   4782
## 2 positive   2006

# Both lexicons have more negative words than positive words, but the bing lexicon has a higher ratio of
# Counting the most frequently appearing words and which sentiment they correspond to (positive or negative)
bing_word_counts <- tidy_combined %>%
  inner_join(get_sentiments("bing")) %>%
  count(word, sentiment, sort = TRUE) %>%
  ungroup()

## Joining, by = "word"
bing_word_counts

## # A tibble: 5,552 x 3
##   word      sentiment      n
##   <chr>    <chr>      <int>
## 1 trump    positive   23953
## 2 like     positive   14612
## 3 well     positive    8250
## 4 right    positive    7530
## 5 good     positive    5802
## 6 work     positive    5544
## 7 support  positive    5504
## 8 free     positive    4327
## 9 great    positive    4007
## 10 strong  positive    3862
## # ... with 5,542 more rows

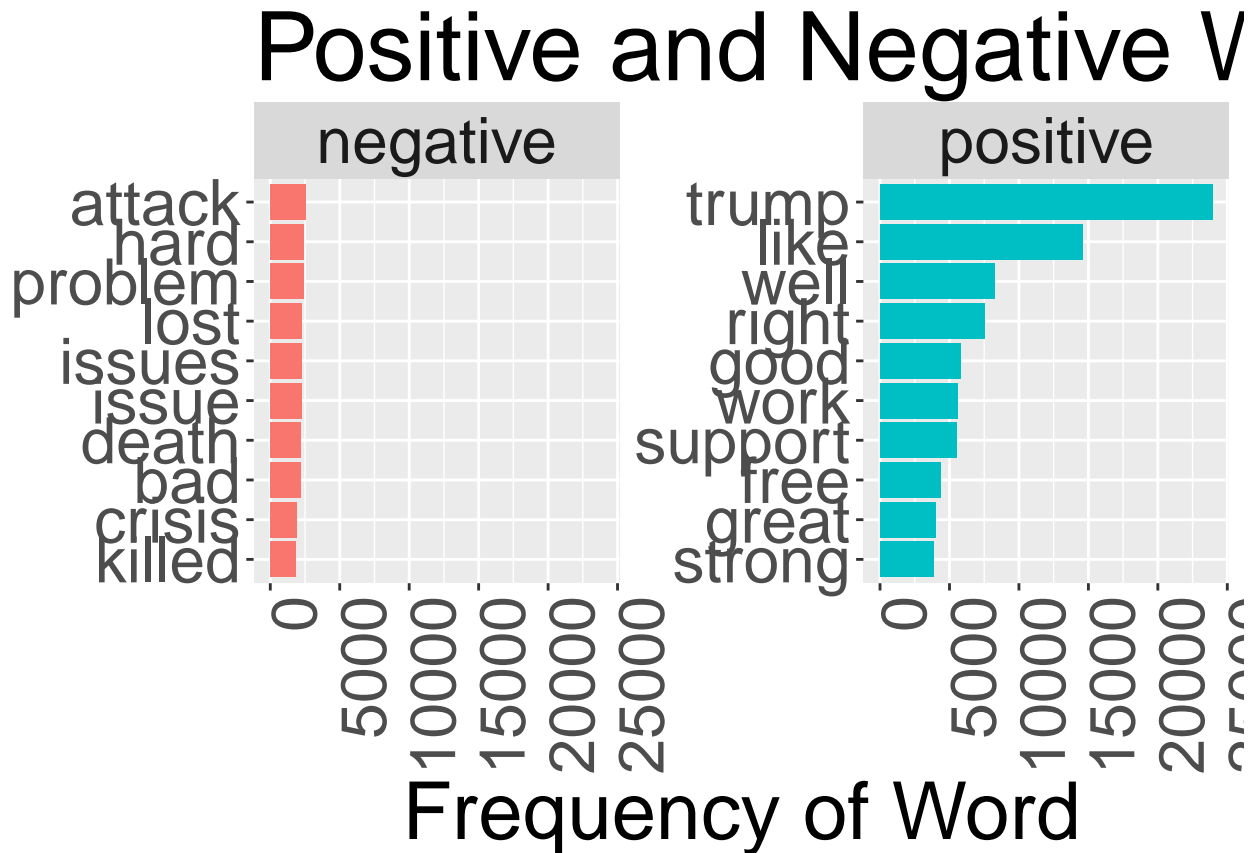
bing_word_counts %>%
  group_by(sentiment) %>%
  top_n(10) %>%
  ungroup() %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(word, n, fill = sentiment)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~sentiment, scales = "free_y") +
  ggtitle("Positive and Negative Word Frequency") +
  labs(y = "Frequency of Word",
       x = NULL) +

```



```
theme(text = element_text(size=30),
      axis.text.x = element_text(angle=90, hjust=1)) +
coord_flip()
```

```
## Selecting by n
```



```
# Wordcloud with most frqeuently appearing words
tidy_combined %>%
  anti_join(stop_words) %>%
  count(word) %>%
  with(wordcloud(words = word, freq = n, max.words = 100, min.freq = 1, random.order=FALSE, rot.per = 0
```

```
## Joining, by = "word"
```

```
## Warning in strwidth(words[i], cex = size[i], ...): conversion failure on
## 'it's' in 'mbcsToSbcs': dot substituted for <e2>
```

```
## Warning in strwidth(words[i], cex = size[i], ...): conversion failure on
## 'it's' in 'mbcsToSbcs': dot substituted for <80>
```

```
## Warning in strwidth(words[i], cex = size[i], ...): conversion failure on
## 'it's' in 'mbcsToSbcs': dot substituted for <99>
```

```
## Warning in text.default(x1, y1, words[i], cex = size[i], offset = 0, srt =
## rotWord * : conversion failure on 'it's' in 'mbcsToSbcs': dot substituted
## for <e2>
```

```
## Warning in text.default(x1, y1, words[i], cex = size[i], offset = 0, srt =
## rotWord * : conversion failure on 'it's' in 'mbcsToSbcs': dot substituted
## for <80>
```

```

## Warning in text.default(x1, y1, words[i], cex = size[i], offset = 0, srt =
## rotWord * : conversion failure on 'it's' in 'mbcsToSbcs': dot substituted
## for <99>

## Warning in text.default(x1, y1, words[i], cex = size[i], offset = 0, srt =
## rotWord * : font metrics unknown for Unicode character U+2019

## Warning in strwidth(words[i], cex = size[i], ...): conversion failure on
## 'don't' in 'mbcsToSbcs': dot substituted for <e2>

## Warning in strwidth(words[i], cex = size[i], ...): conversion failure on
## 'don't' in 'mbcsToSbcs': dot substituted for <80>

## Warning in strwidth(words[i], cex = size[i], ...): conversion failure on
## 'don't' in 'mbcsToSbcs': dot substituted for <99>

## Warning in text.default(x1, y1, words[i], cex = size[i], offset = 0, srt =
## rotWord * : conversion failure on 'don't' in 'mbcsToSbcs': dot substituted
## for <e2>

## Warning in text.default(x1, y1, words[i], cex = size[i], offset = 0, srt =
## rotWord * : conversion failure on 'don't' in 'mbcsToSbcs': dot substituted
## for <80>

## Warning in text.default(x1, y1, words[i], cex = size[i], offset = 0, srt =
## rotWord * : conversion failure on 'don't' in 'mbcsToSbcs': dot substituted
## for <99>

## Warning in text.default(x1, y1, words[i], cex = size[i], offset = 0, srt =
## rotWord * : font metrics unknown for Unicode character U+2019

## Warning in wordcloud(words = word, freq = n, max.words = 100, min.freq =
## 1, : presidential could not be fit on page. It will not be plotted.

## Warning in wordcloud(words = word, freq = n, max.words = 100, min.freq =
## 1, : washington could not be fit on page. It will not be plotted.

## Warning in wordcloud(words = word, freq = n, max.words = 100, min.freq =
## 1, : international could not be fit on page. It will not be plotted.

## Warning in strwidth(words[i], cex = size[i], ...): conversion failure on
## ' ' in 'mbcsToSbcs': dot substituted for <d0>

## Warning in strwidth(words[i], cex = size[i], ...): conversion failure on
## ' ' in 'mbcsToSbcs': dot substituted for <b2>

## Warning in text.default(x1, y1, words[i], cex = size[i], offset = 0, srt =
## rotWord * : conversion failure on ' ' in 'mbcsToSbcs': dot substituted for
## <d0>

## Warning in text.default(x1, y1, words[i], cex = size[i], offset = 0, srt =
## rotWord * : conversion failure on ' ' in 'mbcsToSbcs': dot substituted for
## <b2>

## Warning in text.default(x1, y1, words[i], cex = size[i], offset = 0, srt =
## rotWord * : font metrics unknown for Unicode character U+0432

## Warning in wordcloud(words = word, freq = n, max.words = 100, min.freq =
## 1, : investigation could not be fit on page. It will not be plotted.

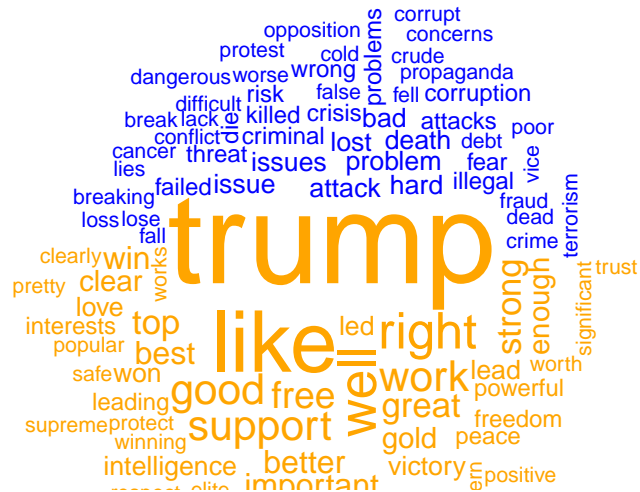
## Warning in wordcloud(words = word, freq = n, max.words = 100, min.freq =
## 1, : democratic could not be fit on page. It will not be plotted.

```

[illegible]

```
# Wordcloud faceted into positive and negative with color (blue corresponds to a negative sentiment whi
tidy_combined %>%
  inner_join(get_sentiments("bing")) %>%
  count(word, sentiment, sort = TRUE) %>%
  acast(word ~ sentiment, value.var = "n", fill = 0) %>%
  comparison.cloud(colors = c("blue", "orange"),
    max.words = 100)
```

```
## Joining, by = "word"
```



```
# Now, it is time to start the machine learning aspect of the project.
```

```
# Using the AFINN lexicon to append the sentiment score of each word to a new dataset called tidy_combi
tidy_combined_a <- tidy_combined %>%
  inner_join(get_sentiments("afinn"))
```

```
## Joining, by = "word"
```

```
# Categorize article as positive or negative overall based on the average of the AFINN score of the wor
tidy_combined_final <- tidy_combined_a %>%
  select(uuid, score, binary_type) %>%
  group_by(uuid) %>%
  summarise(n_words = n(), avgscore = sum(score) / n_words,
    type = mean(binary_type),
    positive_score = sum(score[score > 0]),
    negative_score = sum(score[score < 0]),
    n_positive = sum(score > 0),
    n_negative = sum(score < 0)
  ) %>%
  mutate(articlesent = ifelse(avgscore < 0, "Negative", "Positive")) %>%
  mutate(txt_type = as.factor(type)) %>%
  select(-type)
tidy_combined_final
```

```
## # A tibble: 16,693 x 9
```

```
##   uuid  n_words avgscore positive_score negative_score n_positive
##   <chr>   <int>   <dbl>         <int>         <int>         <int>
## 1 0005~    21    0.286            19            -13            13
```

```
## 2 0020~      24  -0.667          12          -28          7
## 3 0021~      87   0.379         109          -76         49
## 4 002d~      88   0.261          99          -76         50
## 5 0033~       9    0            8            -8          5
## 6 0033~      58  -0.759          36          -80         20
## 7 0037~      14   0.714          16            -6          8
## 8 0038~      30  -0.667          20          -40          9
## 9 003d~      10   0.7           14            -7          7
## 10 0048~     58   0.172          50          -40         34
## # ... with 16,683 more rows, and 3 more variables: n_negative <int>,
## #   articlesent <chr>, txt_type <fct>
```

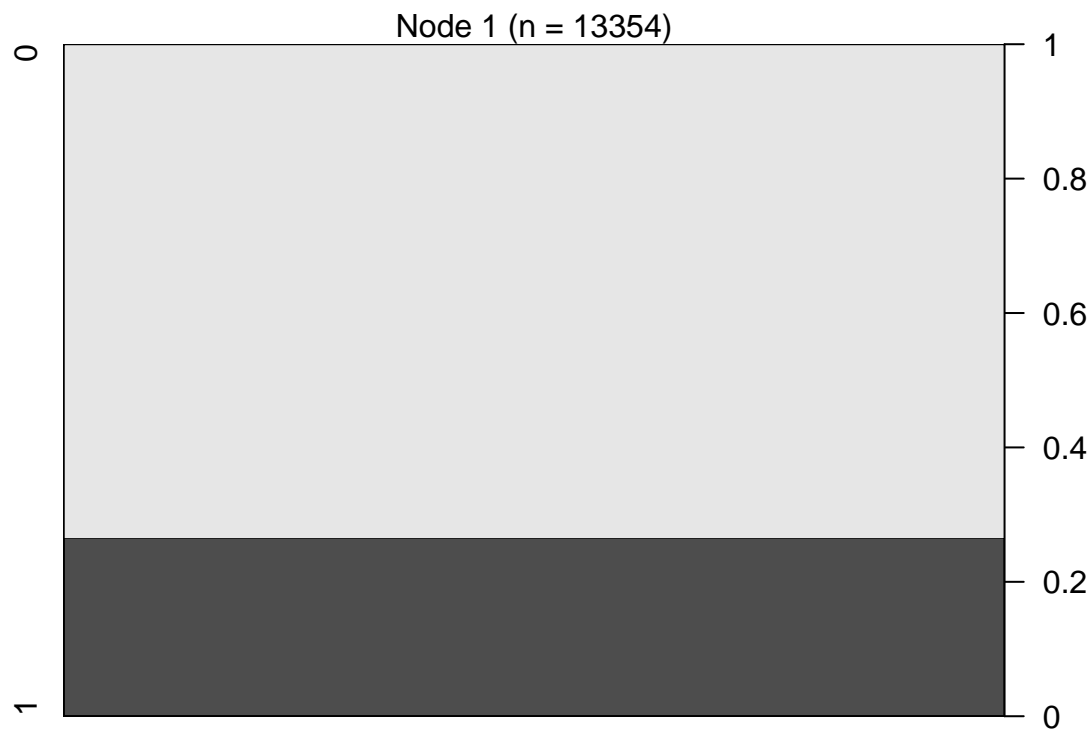
```
tidy_combined_final %>%
  filter(txt_type == 0) %>%
  summarise(n_negative = n())
```

```
## # A tibble: 1 x 1
##   n_negative
##       <int>
## 1      12248
```

```
# Decision tree training process
```

```
n <- nrow(tidy_combined_final)
train_id <- sample(1:n, size = round(n * 0.8))
train <- tidy_combined_final[train_id,]
test <- tidy_combined_final[-train_id,]
```

```
tree <- rpart(txt_type ~ avgscore + n_words + n_positive + n_negative + negative_score + positive_score)
plot(as.party(tree))
```



```

tree

## n= 13354
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 13354 3550 0 (0.7341620 0.2658380) *

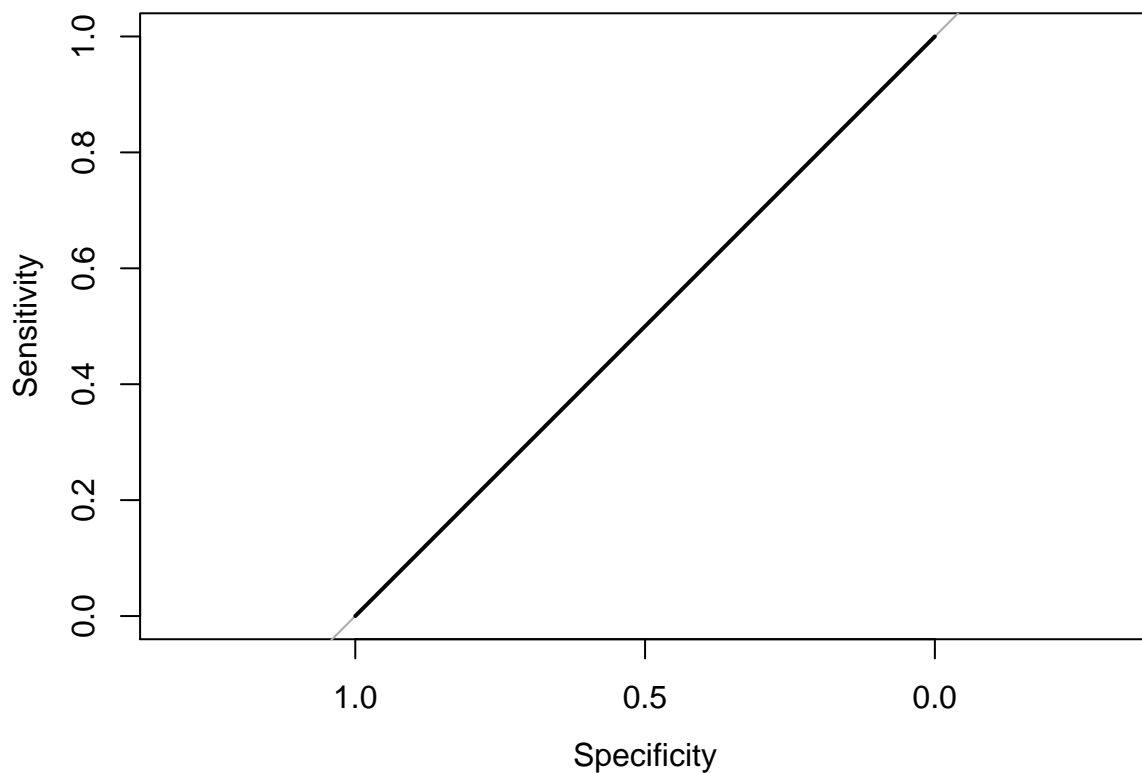
saveRDS(tree, file = "tree.rds")
saveRDS(train, file = "train.rds")
prediction <- predict(tree, test)

test <- test %>%
  mutate(prediction = prediction[1])
roc_obj <- roc(test$txt_type, test$prediction)
auc(roc_obj)

## Area under the curve: 0.5

plot(roc_obj)

```



```

# Based on this tree, we can see that none of the predictors (average score, number of words, number of

# Why is this true? Below, we will do some exploration using visualizations to display the poor relation

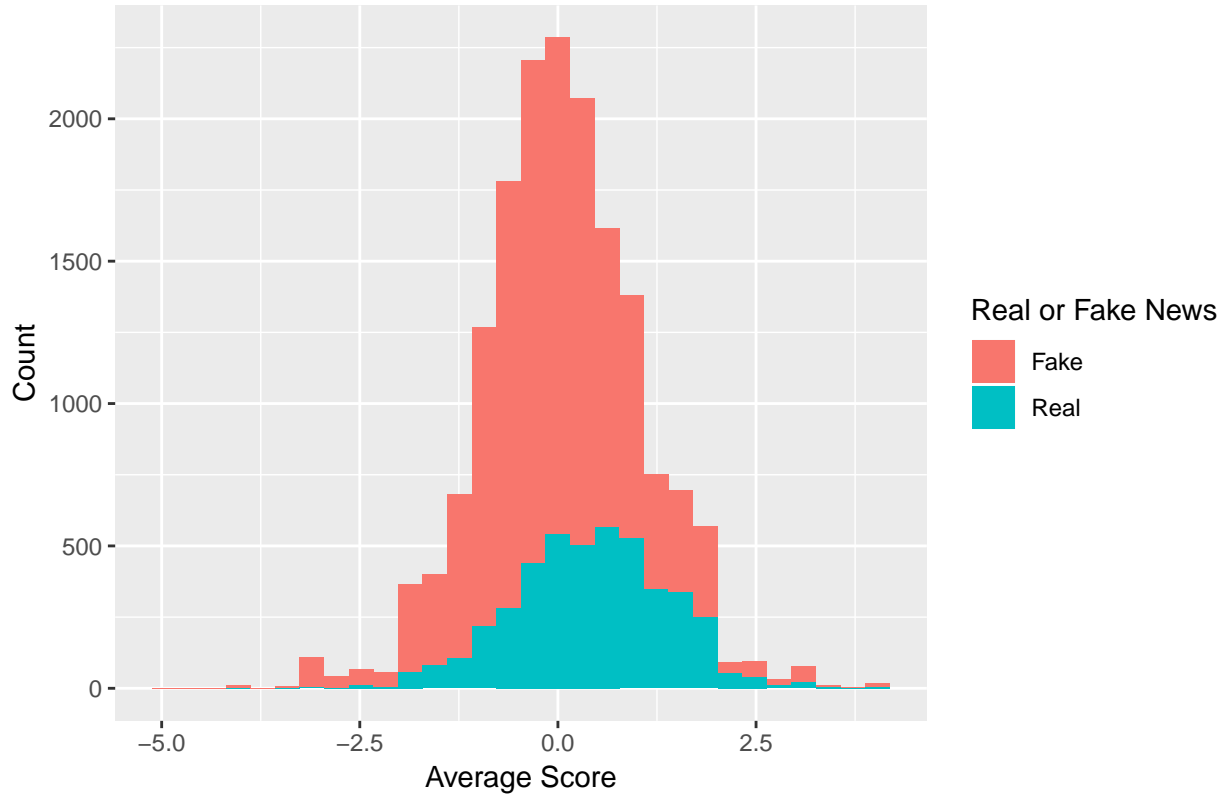
# Histogram of Average Sentiment Score by News Type (Real and Fake)
ggplot(tidy_combined_final, aes(x = avgscore, fill = txt_type)) +
  geom_histogram() +
  xlab("Average Score") +
  ylab("Count") +

```

```
ggtitle("Histogram of Average Sentiment Score by News Type (Real and Fake)") +
scale_fill_discrete(name = "Real or Fake News", labels = c("Fake", "Real"))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Histogram of Average Sentiment Score by News Type (Real and Fake)

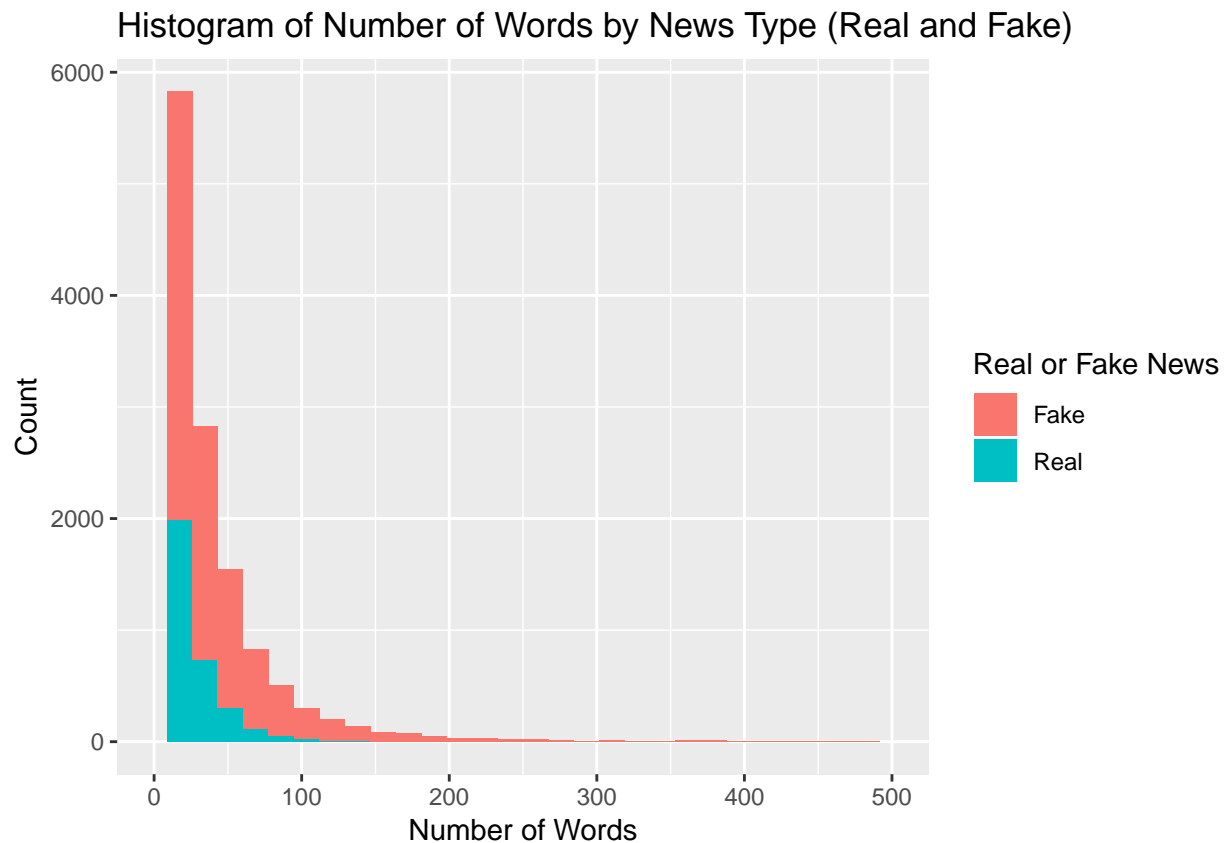


```
# Histogram of Number of Words by News Type (Real and Fake)
ggplot(tidy_combined_final, aes(x = n_words, fill = txt_type)) +
  geom_histogram() +
  xlim(0, 500) +
  xlab("Number of Words") +
  ylab("Count") +
  ggtitle("Histogram of Number of Words by News Type (Real and Fake)") +
  scale_fill_discrete(name = "Real or Fake News", labels = c("Fake", "Real"))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 25 rows containing non-finite values (stat_bin).
```

```
## Warning: Removed 4 rows containing missing values (geom_bar).
```



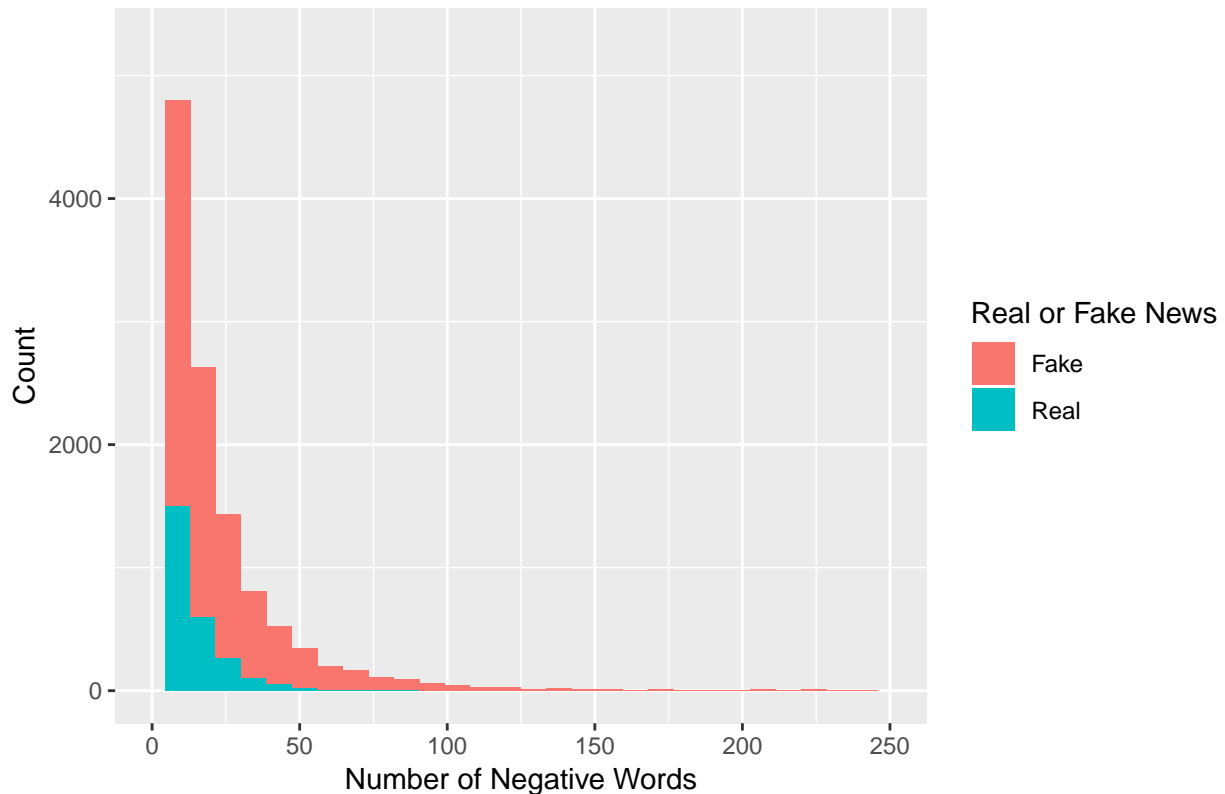
```
# Histogram of Number of Negative Words by News Type (Real and Fake)
ggplot(tidy_combined_final, aes(x = n_negative, fill = txt_type)) +
  geom_histogram() +
  xlim(0, 250) +
  xlab("Number of Negative Words") +
  ylab("Count") +
  ggtitle("Histogram of Number of Negative Words by News Type (Real and Fake)") +
  scale_fill_discrete(name = "Real or Fake News", labels = c("Fake", "Real"))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 27 rows containing non-finite values (stat_bin).
```

```
## Warning: Removed 4 rows containing missing values (geom_bar).
```


Histogram of Number of Negative Words by News Type (Real and Fake)



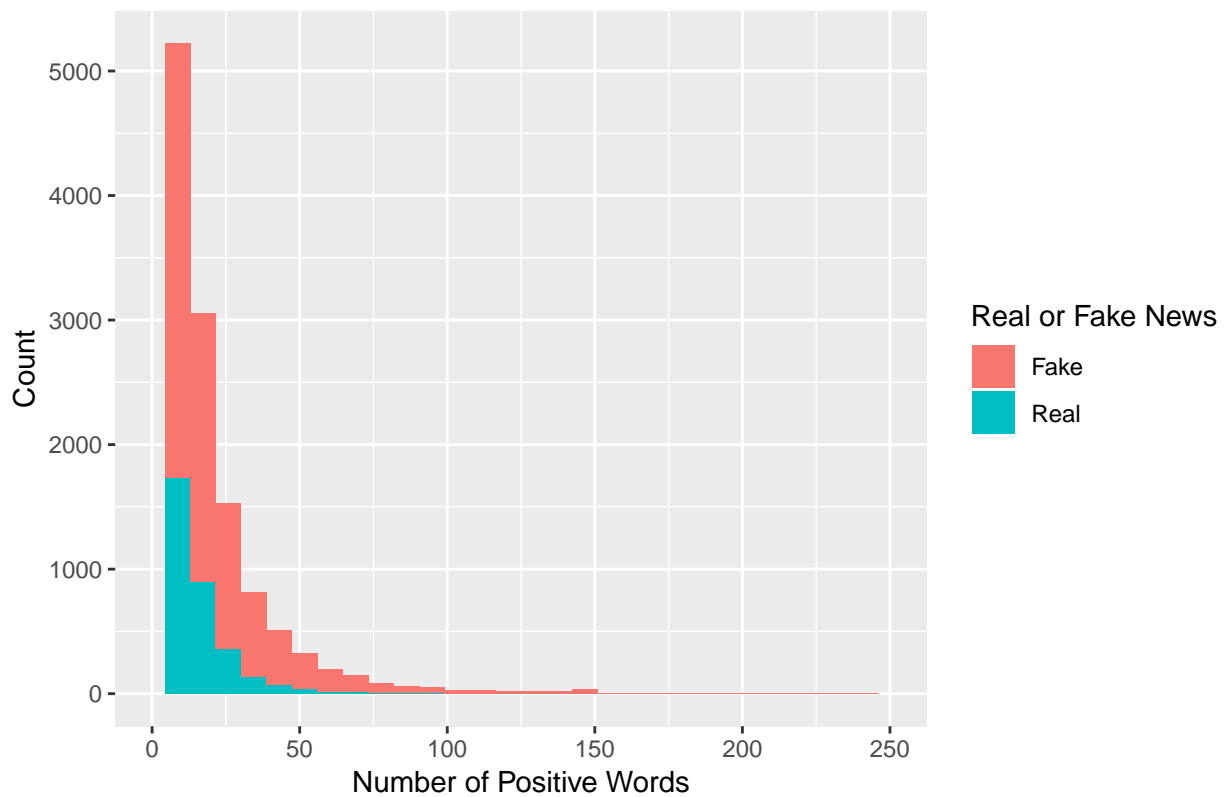
```
# Histogram of Number of Positive Words by News Type (Real and Fake
ggplot(tidy_combined_final, aes(x = n_positive, fill = txt_type)) +
  geom_histogram() +
  xlim(0, 250) +
  xlab("Number of Positive Words") +
  ylab("Count") +
  ggtitle("Histogram of Number of Positive Words by News Type (Real and Fake)") +
  scale_fill_discrete(name = "Real or Fake News", labels = c("Fake", "Real"))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 27 rows containing non-finite values (stat_bin).
```

```
## Warning: Removed 4 rows containing missing values (geom_bar).
```

Histogram of Number of Positive Words by News Type (Real and Fake)



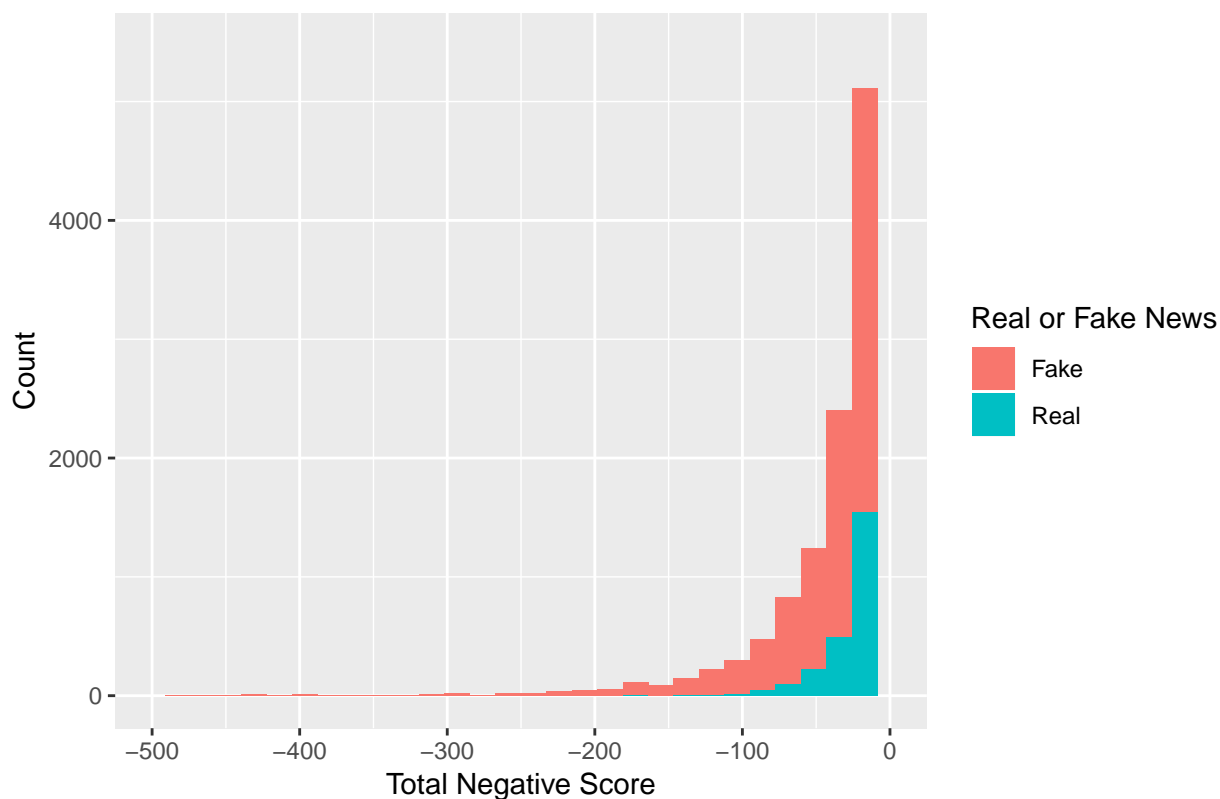
```
# Histogram of Total Negative Score by News Type (Real and Fake)
ggplot(tidy_combined_final, aes(x = negative_score, fill = txt_type)) +
  geom_histogram() +
  xlim(-500, 0) +
  xlab("Total Negative Score") +
  ylab("Count") +
  ggtitle("Histogram of Total Negative Score by News Type (Real and Fake)") +
  scale_fill_discrete(name = "Real or Fake News", labels = c("Fake", "Real"))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 25 rows containing non-finite values (stat_bin).
```

```
## Warning: Removed 4 rows containing missing values (geom_bar).
```

Histogram of Total Negative Score by News Type (Real and Fake)



```
# Histogram of Total Positive Score by News Type (Real and Fake)
ggplot(tidy_combined_final, aes(x = positive_score, fill = txt_type)) +
  geom_histogram() +
  xlim(0, 500) +
  xlab("Total Positive Score") +
  ylab("Count") +
  ggtitle("Histogram of Total Positive Score by News Type (Real and Fake)") +
  scale_fill_discrete(name = "Real or Fake News", labels = c("Fake", "Real"))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 21 rows containing non-finite values (stat_bin).
```

```
## Warning: Removed 4 rows containing missing values (geom_bar).
```

Histogram of Total Positive Score by News Type (Real and Fake)

