

Detecting Fake News Using Sentiment Analysis

Andrea Boskovic and Peter Cho

12/18/2018

```
# Loading necessary packages
```

```
knitr::opts_chunk$set(echo = TRUE)
```

```
# Tidying data packages
```

```
library(ggplot2) #dynamic graphics
```

```
library(tidyr) #data tidying
```

```
library(readr) #reading in csv file
```

```
library(dplyr) #manipulating datasets
```

```
# Text Data Extraction and Manipulation
```

```
library(tm) #text mining
```

```
library(tidytext) #text tidying
```

```
library(wordcloud) #to make wordcloud
```

```
library(reshape2) #to make faceted wordcloud
```

```
# Decision Tree Modeling Packages
```

```
library(rpart) #tree modeling for classification
```

```
library(partykit) #tree modeling
```

```
library(pROC) #area under the curve
```

```
#Setting the seed for functions involving randomness
```

```
set.seed(123)
```

```
# Loading the dataset
```

```
fake <- read_csv("fake.csv") #all fake
```

```
real <- read_csv("Articles.csv") #all real
```

```
new_ds <- read_csv("data.csv") #combination of real and fake
```

```
fake_type <- c("fake", "satire", "bias", "bs", "conspiracy", "state", "junksci", "hate")
```

```
real_type <- c("sports", "business")
```

```
# Merging the datasets and removing unnecessary columns
```

```
real <- real %>%
```

```
  mutate(binary_type = ifelse(NewsType %in% fake_type, 0, 1)) #now fake = 0 and real = 1
```

```
fake <- fake %>%
```

```
  mutate(binary_type = ifelse(type %in% fake_type, 0, 1)) #now fake = 0 and real = 1
```

```
new_ds <- new_ds %>%
```

```
  filter(Label == 1)
```

```
real <- full_join(real, new_ds, by = c("Heading" = "Headline", "Article" = "Body", "binary_type" = "Label"))
```

```
real <- real %>%
```

```
  mutate(id = as.character(seq(1:4564))) %>%
```

```
  mutate(realtype = "real")
```

```
# Making a combined dataset with both fake and real articles and selecting only for the uuid (unique id)
```

```
combined <- full_join(fake, real, by = c("text" = "Article", "title" = "Heading", "uuid" = "id", "binary_type" = "Label"))
```

```
select(uuid, binary_type, type, title, text)
```

```
# Making a tidy dataset where we have the the words in their own column for facilitated data analysis a
tidy_combined <- combined %>%
  unnest_tokens(word, text)
head(tidy_combined)
```

```
## # A tibble: 6 x 5
##   uuid                binary_type type title                word
##   <chr>                <dbl> <chr> <chr>                <chr>
## 1 6a175f46bcd24d39b3e9~      0 bias Muslims BUSTED: They Stol~ print
## 2 6a175f46bcd24d39b3e9~      0 bias Muslims BUSTED: They Stol~ they
## 3 6a175f46bcd24d39b3e9~      0 bias Muslims BUSTED: They Stol~ shou~
## 4 6a175f46bcd24d39b3e9~      0 bias Muslims BUSTED: They Stol~ pay
## 5 6a175f46bcd24d39b3e9~      0 bias Muslims BUSTED: They Stol~ all
## 6 6a175f46bcd24d39b3e9~      0 bias Muslims BUSTED: They Stol~ the
```

```
# Basic Data Exploration:
```

```
# This allows us to see how many observations are in each type of fake news.
combined %>%
  group_by(type) %>%
  summarize(n = n())
```

```
## # A tibble: 9 x 2
##   type                n
##   <chr>              <int>
## 1 bias              443
## 2 bs              11492
## 3 conspiracy        430
## 4 fake              19
## 5 hate             246
## 6 junksci          102
## 7 real            4564
## 8 satire           146
## 9 state            121
```

```
typetotals <- combined %>%
  group_by(type) %>%
  summarize(n = n())
typetotals
```

```
## # A tibble: 9 x 2
##   type                n
##   <chr>              <int>
## 1 bias              443
## 2 bs              11492
## 3 conspiracy        430
## 4 fake              19
## 5 hate             246
## 6 junksci          102
## 7 real            4564
## 8 satire           146
## 9 state            121
```

```
# What are the most common words for each basic emotion?
# We will use the nrc lexicon to categorize each documented word into on of the basic human emotions ca
```

```
# Anger
nrc_anger <- get_sentiments("nrc") %>%
  filter(sentiment == "anger")

tidy_combined %>%
  inner_join(nrc_anger) %>%
  count(word, sort = TRUE) %>%
  head(n = 10)
```

```
## # A tibble: 10 x 2
##   word      n
##   <chr>    <int>
## 1 vote      4969
## 2 money      4835
## 3 force      3189
## 4 court      2721
## 5 attack      2548
## 6 defense      2242
## 7 death      2176
## 8 bad         2175
## 9 politics    2058
## 10 fight      2054
```

```
# Fear
nrc_fear <- get_sentiments("nrc") %>%
  filter(sentiment == "fear")

tidy_combined %>%
  inner_join(nrc_fear) %>%
  count(word, sort = TRUE) %>%
  head(n = 10)
```

```
## # A tibble: 10 x 2
##   word      n
##   <chr>    <int>
## 1 government 11656
## 2 war        9845
## 3 military    5880
## 4 police      4902
## 5 change      4442
## 6 case        4177
## 7 force       3189
## 8 court       2721
## 9 attack      2548
## 10 problem    2381
```

```
# Anticipation
nrc_anticipation <- get_sentiments("nrc") %>%
  filter(sentiment == "anticipation")

tidy_combined %>%
  inner_join(nrc_anticipation) %>%
  count(word, sort = TRUE) %>%
  head(n = 10)
```

```
## # A tibble: 10 x 2
```

```
##      word      n
##      <chr>    <int>
## 1 time      14159
## 2 white     6547
## 3 public    6039
## 4 good      5802
## 5 long      5706
## 6 vote      4969
## 7 money     4835
## 8 investigation 3968
## 9 top       3822
## 10 continue 3439
```

```
# Trust
nrc_trust <- get_sentiments("nrc") %>%
  filter(sentiment == "trust")

tidy_combined %>%
  inner_join(nrc_trust) %>%
  count(word, sort = TRUE) %>%
  head(n = 10)
```

```
## # A tibble: 10 x 2
##      word      n
##      <chr>    <int>
## 1 president 12344
## 2 united    7803
## 3 white     6547
## 4 good      5802
## 5 law       5181
## 6 system    5088
## 7 vote      4969
## 8 police    4902
## 9 money     4835
## 10 fact     4673
```

```
# Surprise
nrc_surprise <- get_sentiments("nrc") %>%
  filter(sentiment == "surprise")

tidy_combined %>%
  inner_join(nrc_surprise) %>%
  count(word, sort = TRUE) %>%
  head(n = 10)
```

```
## # A tibble: 10 x 2
##      word      n
##      <chr> <int>
## 1 trump  23953
## 2 good   5802
## 3 vote   4969
## 4 money  4835
## 5 deal   2802
## 6 death  2176
## 7 leave  2080
## 8 hope   1902
```

```
## 9 young 1859
## 10 shot 1604
```

Sadness

```
nrc_sadness <- get_sentiments("nrc") %>%
  filter(sentiment == "sadness")
```

```
tidy_combined %>%
  inner_join(nrc_sadness) %>%
  count(word, sort = TRUE) %>%
  head(n = 10)
```

```
## # A tibble: 10 x 2
```

```
##   word      n
##   <chr>   <int>
## 1 vote    4969
## 2 black   4196
## 3 case    4177
## 4 problem 2381
## 5 lost    2260
## 6 tax      2211
## 7 death   2176
## 8 bad      2175
## 9 leave   2080
## 10 violence 1955
```

Joy

```
nrc_joy <- get_sentiments("nrc") %>%
  filter(sentiment == "joy")
```

```
tidy_combined %>%
  inner_join(nrc_joy) %>%
  count(word, sort = TRUE) %>%
  head(n = 10)
```

```
## # A tibble: 10 x 2
```

```
##   word      n
##   <chr> <int>
## 1 white  6547
## 2 good   5802
## 3 vote   4969
## 4 money  4835
## 5 found  4192
## 6 share  3090
## 7 deal   2802
## 8 food   2756
## 9 pay    2339
## 10 true  2234
```

Disgust

```
nrc_disgust <- get_sentiments("nrc") %>%
  filter(sentiment == "disgust")
```

```
tidy_combined %>%
  inner_join(nrc_disgust) %>%
  count(word, sort = TRUE) %>%
```

```

head(n = 10)

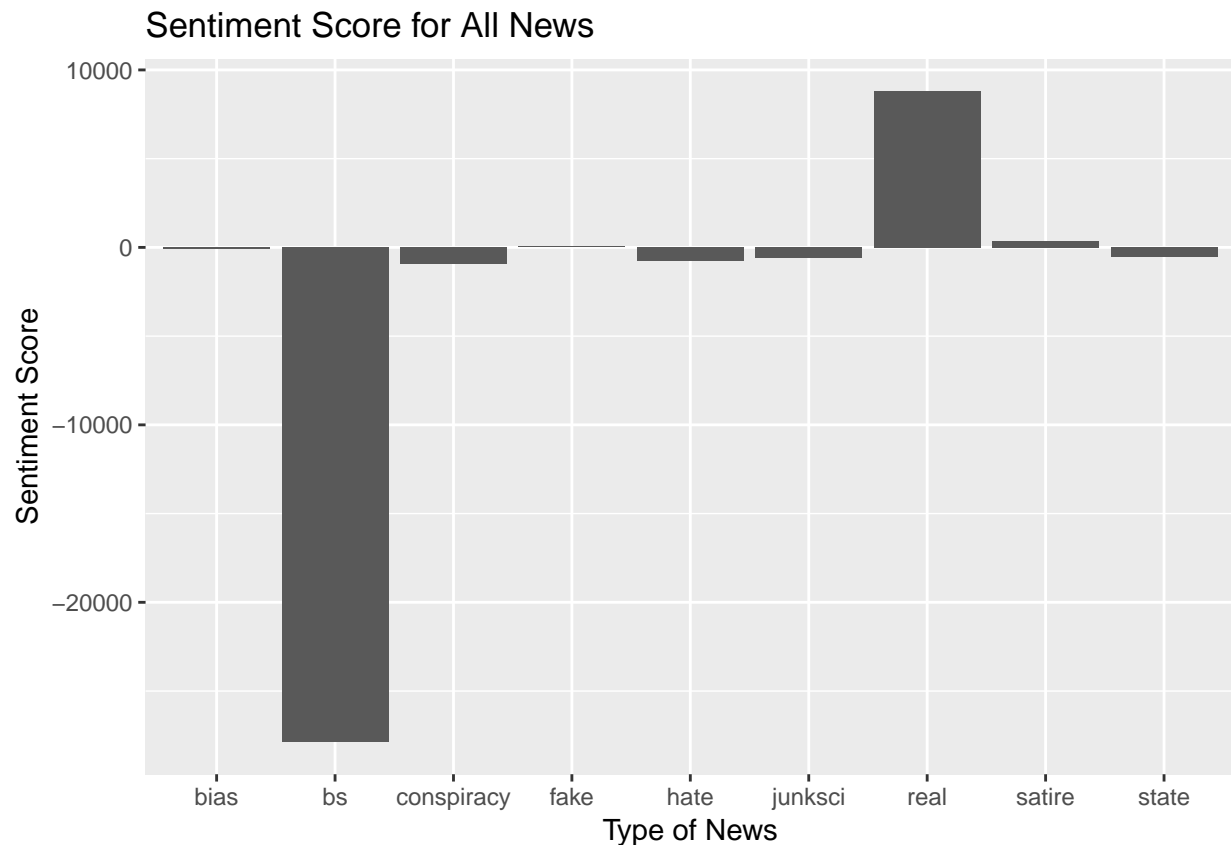
## # A tibble: 10 x 2
##   word      n
##   <chr>   <int>
## 1 john    3108
## 2 congress 2473
## 3 death    2176
## 4 bad      2175
## 5 criminal 1805
## 6 illegal  1756
## 7 powerful 1611
## 8 corruption 1571
## 9 finally  1442
## 10 remains 1244

# Find net sentiment for each type of fake news documented in the dataset using the bing lexicon. The b
# Note that some types, such as bs (> 400,000), have more corresponding observations than other types, .
combined_sentiment <- tidy_combined %>%
  inner_join(get_sentiments("bing")) %>%
  count(type, sentiment) %>%
  spread(sentiment, n, fill = 0) %>%
  mutate(sentiment = positive - negative)
combined_sentiment

## # A tibble: 9 x 4
##   type      negative positive sentiment
##   <chr>      <dbl>   <dbl>   <dbl>
## 1 bias        5422     5322     -100
## 2 bs       247391    219536   -27855
## 3 conspiracy   4805     3851    -954
## 4 fake         148      199      51
## 5 hate        8765     7998   -767
## 6 junksci     3070     2469   -601
## 7 real       45896    54690    8794
## 8 satire      1148     1487     339
## 9 state       1215      704    -511

# Plot of the sentiment score for each type of news
ggplot(combined_sentiment, aes(x = type, y = sentiment)) + geom_col() + labs(title = "Sentiment Score f

```



```
# We can also get the sentiment score on a scale of -5 to 5 from the AFINN lexicon. The AFINN lexicon has
afinn <- tidy_combined %>%
  inner_join(get_sentiments("afinn")) %>%
  group_by(type) %>%
  summarise(sentiment = sum(score)) %>%
  mutate(method = "AFINN")
```

```
## Joining, by = "word"
```

```
head(afinn)
```

```
## # A tibble: 6 x 3
##   type      sentiment method
##   <chr>      <int> <chr>
## 1 bias      -1507 AFINN
## 2 bs       -62021 AFINN
## 3 conspiracy -1846 AFINN
## 4 fake        108 AFINN
## 5 hate      -1625 AFINN
## 6 junksci     41 AFINN
```

```
# It may be useful to investigate the basic contents of the lexicons.
```

```
# Positive and negative words in nrc lexicon
get_sentiments("nrc") %>%
  filter(sentiment %in% c("positive",
                        "negative")) %>%
  count(sentiment)
```

```
## # A tibble: 2 x 2
##   sentiment      n
##   <chr>      <int>
## 1 negative   3324
## 2 positive   2312
```

```
# Positive and negative words in Bing lexicon
get_sentiments("bing") %>%
  count(sentiment)
```

```
## # A tibble: 2 x 2
##   sentiment      n
##   <chr>      <int>
## 1 negative   4782
## 2 positive   2006
```


Both lexicons have more negative words than positive words, but the Bing lexicon has a higher ratio of negative words.

```
# Counting the most frequently appearing words and which sentiment they correspond to (positive or negative)
bing_word_counts <- tidy_combined %>%
  inner_join(get_sentiments("bing")) %>%
  count(word, sentiment, sort = TRUE) %>%
  ungroup()
head(bing_word_counts)
```

```
## # A tibble: 6 x 3
##   word sentiment      n
##   <chr> <chr>      <int>
## 1 trump positive  23953
## 2 like  positive  14612
## 3 well  positive   8250
## 4 right positive   7530
## 5 good  positive   5802
## 6 work  positive   5544
```

```
bing_word_counts %>%
  group_by(sentiment) %>%
  top_n(10) %>%
  ungroup() %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(word, n, fill = sentiment)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~sentiment, scales = "free_y") +
  ggtitle("Positive and Negative Word Frequency") +
  labs(y = "Frequency of Word",
       x = NULL) +
  theme(text = element_text(size=30),
        axis.text.x = element_text(angle=90, hjust=1)) +
  coord_flip()
```


Positive and Negative V



Frequency of Word

```
# Wordcloud with most frquently appearing words
tidy_combined %>%
  anti_join(stop_words) %>%
  count(word) %>%
  with(wordcloud(words = word, freq = n, max.words = 100, min.freq = 1, random.order=FALSE, rot.per = 0
```



```

# Categorize article as positive or negative overall based on the average of the AFINN score of the words
tidy_combined_final <- tidy_combined_a %>%
  select(uuid, score, binary_type) %>%
  group_by(uuid) %>%
  summarise(n_words = n(), avgscore = sum(score) / n_words,
            type = mean(binary_type),
            positive_score = sum(score[score > 0]),
            negative_score = sum(score[score < 0]),
            n_positive = sum(score > 0),
            n_negative = sum(score < 0)
            ) %>%
  mutate(articlesent = ifelse(avgscore < 0, "Negative", "Positive")) %>%
  mutate(txt_type = as.factor(type)) %>%
  select(-type)
head(tidy_combined_final)

```

```

## # A tibble: 6 x 9
##   uuid  n_words avgscore positive_score negative_score n_positive
##   <chr>   <int>   <dbl>         <int>         <int>         <int>
## 1 0005~     21    0.286             19             -13             13
## 2 0020~     24   -0.667             12             -28              7
## 3 0021~     87    0.379            109             -76             49
## 4 002d~     88    0.261             99             -76             50
## 5 0033~      9     0              8              -8              5
## 6 0033~     58   -0.759             36             -80             20
## # ... with 3 more variables: n_negative <int>, articlesent <chr>,
## #   txt_type <fct>

```

```

tidy_combined_final %>%
  filter(txt_type == 0) %>%
  summarise(n_negative = n())

```

```

## # A tibble: 1 x 1
##   n_negative
##   <int>
## 1     12248

```

```

# Decision tree training process

```

```


n <- nrow(tidy_combined_final)
train_id <- sample(1:n, size = round(n * 0.8))
train <- tidy_combined_final[train_id,]
test <- tidy_combined_final[-train_id,]

```

```

tree <- rpart(txt_type ~ avgscore + n_words + n_positive + n_negative + negative_score + positive_score)
plot(as.party(tree))

```



```
tree
```


```
## n= 13354
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 13354 3550 0 (0.7341620 0.2658380) *
```

```
prediction <- predict(tree, test)
```

```
test <- test %>%
  mutate(prediction = prediction[1])
roc_obj <- roc(test$txt_type, test$prediction)
auc(roc_obj)
```

```
## Area under the curve: 0.5
```

```
plot(roc_obj)
```




Based on this tree, we can see that none of the predictors (average score, number of words, number of

Why is this true? Below, we will do some exploration using visualizations to display the poor relation

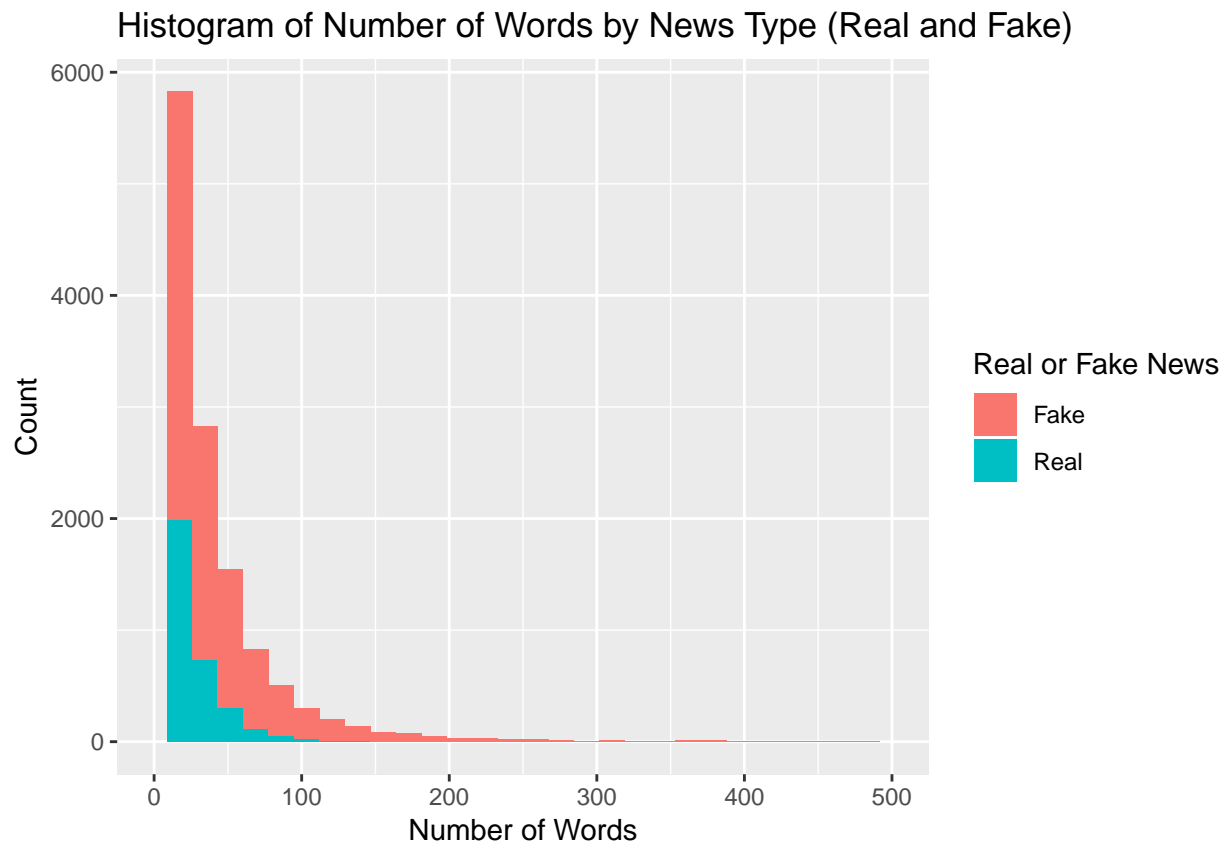
Histogram of Average Sentiment Score by News Type (Real and Fake)

```
ggplot(tidy_combined_final, aes(x = avgscore, fill = txt_type)) +
  geom_histogram() +
  xlab("Average Score") +
  ylab("Count") +
  ggtitle("Histogram of Average Sentiment Score by News Type (Real and Fake)") +
  scale_fill_discrete(name = "Real or Fake News", labels = c("Fake", "Real"))
```

Histogram of Average Sentiment Score by News Type (Real and Fake)




```
# Histogram of Number of Words by News Type (Real and Fake)
ggplot(tidy_combined_final, aes(x = n_words, fill = txt_type)) +
  geom_histogram() +
  xlim(0, 500) +
  xlab("Number of Words") +
  ylab("Count") +
  ggtitle("Histogram of Number of Words by News Type (Real and Fake)") +
  scale_fill_discrete(name = "Real or Fake News", labels = c("Fake", "Real"))
```




```
# Histogram of Number of Negative Words by News Type (Real and Fake)
ggplot(tidy_combined_final, aes(x = n_negative, fill = txt_type)) +
  geom_histogram() +
  xlim(0, 250) +
  xlab("Number of Negative Words") +
  ylab("Count") +
  ggtitle("Histogram of Number of Negative Words by News Type (Real and Fake)") +
  scale_fill_discrete(name = "Real or Fake News", labels = c("Fake", "Real"))
```

Histogram of Number of Negative Words by News Type (Real and Fake)




```
# Histogram of Number of Positive Words by News Type (Real and Fake)
ggplot(tidy_combined_final, aes(x = n_positive, fill = txt_type)) +
  geom_histogram() +
  xlim(0, 250) +
  xlab("Number of Positive Words") +
  ylab("Count") +
  ggtitle("Histogram of Number of Positive Words by News Type (Real and Fake)") +
  scale_fill_discrete(name = "Real or Fake News", labels = c("Fake", "Real"))
```


Histogram of Number of Positive Words by News Type (Real and Fake)



```
# Histogram of Total Negative Score by News Type (Real and Fake)
ggplot(tidy_combined_final, aes(x = negative_score, fill = txt_type)) +
  geom_histogram() +
  xlim(-500, 0) +
  xlab("Total Negative Score") +
  ylab("Count") +
  ggtitle("Histogram of Total Negative Score by News Type (Real and Fake)") +
  scale_fill_discrete(name = "Real or Fake News", labels = c("Fake", "Real"))
```

Histogram of Total Negative Score by News Type (Real and Fake)



```
# Histogram of Total Positive Score by News Type (Real and Fake)
ggplot(tidy_combined_final, aes(x = positive_score, fill = txt_type)) +
  geom_histogram() +
  xlim(0, 500) +
  xlab("Total Positive Score") +
  ylab("Count") +
  ggtitle("Histogram of Total Positive Score by News Type (Real and Fake)") +
  scale_fill_discrete(name = "Real or Fake News", labels = c("Fake", "Real"))
```

Histogram of Total Positive Score by News Type (Real and Fake)

